

STATS 506 Problem Set #1

Haiming Li

Wine Data

a. Data loading.

```
wine_df <- read.csv('./wine/wine.data', header=FALSE)
names(wine_df) <- c('class', 'alcohol', 'malic_acid', 'ash',
                    'alcalinity_of_ash', 'magnesium', 'total_phenols',
                    'flavanoids', 'nonflavanoid_phenols', 'proanthocyanins',
                    'color_intensity', 'hue', 'od280_od315_ratio', 'proline')
```

b. Class count seems to be consistent.

```
table(wine_df$class)
```

```
1  2  3
59 71 48
```

c. Here are the answers.

1. Correlation is 0.5463642.

```
cor(wine_df$alcohol, wine_df$color_intensity)
```

```
[1] 0.5463642
```

2. It seems class 1 have the highest correlation and class 2 have the lowest.

```
splited_df <- split(wine_df, wine_df$class)
sapply(splited_df, function(df) cor(df$alcohol, df$color_intensity))
```

	1	2	3
	0.4082913	0.2697891	0.3503777

3. The alcohol content of the wine with the highest color intensity is 14.34.

```
wine_df$alcohol[which.max(wine_df$color_intensity)]
```

```
[1] 14.34
```

4. 8.43% of wines had a higher content of proanthocyanins compare to ash.

```
mean(wine_df$proanthocyanins > wine_df$ash)
```

```
[1] 0.08426966
```

d. Here's the table. Notice that here I used a previously defined variable that splits the dataframe based on class.

```
res <- data.frame(
  overall=colMeans(wine_df[, -1]),
  class1=colMeans(splited_df$"1"[, -1]),
  class2=colMeans(splited_df$"2"[, -1]),
  class3=colMeans(splited_df$"3"[, -1])
)
data.frame(t(res))
```

	alcohol	malic_acid	ash	alcalinity_of_ash	magnesium	total_phenols
overall	13.00062	2.336348	2.366517	19.49494	99.74157	2.295112
class1	13.74475	2.010678	2.455593	17.03729	106.33898	2.840169
class2	12.27873	1.932676	2.244789	20.23803	94.54930	2.258873
class3	13.15375	3.333750	2.437083	21.41667	99.31250	1.678750
	flavanoids	nonflavanoid_phenols	proanthocyanins	color_intensity		
overall	2.0292697		0.3618539	1.590899		5.058090
class1	2.9823729		0.2900000	1.899322		5.528305
class2	2.0808451		0.3636620	1.630282		3.086620
class3	0.7814583		0.4475000	1.153542		7.396250
	hue	od280_od315_ratio	proline			
overall	0.9574494	2.611685	746.8933			
class1	1.0620339	3.157797	1115.7119			
class2	1.0562817	2.785352	519.5070			
class3	0.6827083	1.683542	629.8958			

- e. Here we can see all p-values are less than 0.05, thus we can reject the null hypothesis. The phenols level appears to be different across classes.

```
# class 1 v.s. class 2
t.test(splited_df$"1"$total_phenols, splited_df$"2"$total_phenols)
```

Welch Two Sample t-test

```
data: splited_df$"1"$total_phenols and splited_df$"2"$total_phenols
t = 7.4206, df = 119.14, p-value = 1.889e-11
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 0.4261870 0.7364055
sample estimates:
mean of x mean of y
 2.840169  2.258873
```

```
# class 1 v.s. class 3
t.test(splited_df$"1"$total_phenols, splited_df$"3"$total_phenols)
```

Welch Two Sample t-test

```
data: splited_df$"1"$total_phenols and splited_df$"3"$total_phenols
t = 17.12, df = 98.356, p-value < 2.2e-16
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 1.026801 1.296038
sample estimates:
mean of x mean of y
 2.840169  1.678750
```

```
# class 2 v.s. class 3
t.test(splited_df$"2"$total_phenols, splited_df$"3"$total_phenols)
```

Welch Two Sample t-test

```
data: splited_df$"2"$total_phenols and splited_df$"3"$total_phenols
```

```
t = 7.0125, df = 116.91, p-value = 1.622e-10
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 0.4162855 0.7439610
sample estimates:
mean of x mean of y
 2.258873  1.678750
```

AskAManager.org Data

- a. Data loading. (I removed the useless index column)

```
df <- read.csv("AskAManager.csv")[-1]
```

- b. Column cleaning.

```
names(df) <- c('timestamp', 'age', 'industry', 'job_title', 'job_title_context',
               'salary', 'additional_income', 'income_currency', 'other_currency',
               'income_context', 'country', 'state', 'city', 'yoe_total', 'yoe_in_field',
               'highest_education', 'gender', 'race')
```

- c. Filter by income currency.

```
# make necessary columns as factor
df$income_currency <- as.factor(df$income_currency)
print(paste('Obs before:', dim(df)[1]))
```

```
[1] "Obs before: 28062"
```

```
df <- subset(df, income_currency == 'USD')
print(paste('Obs after: ', dim(df)[1]))
```

```
[1] "Obs after: 23374"
```

- d. Remove impossible working experience. My filtering logic that that total year of experience should be at least as much as year of experience in the field. For people in age range 18 ~ 24, it is impossible to gain total yoe beyond 7 years; for people in age range 25 ~ 34, it is impossible to gain total yoe beyond 16 years; for people in age range 35 ~ 44, it is impossible to gain total yoe beyond 26 years; for people in age range 45 ~ 54, it is impossible to gain total yoe beyond 36 years.

```

# convert necessary columns to factor
yoe_levels <- c("1 year or less", "2 - 4 years", "5-7 years", "8 - 10 years",
               "11 - 20 years", "21 - 30 years", "31 - 40 years", "41 years or more")
df$age <- as.factor(df$age)
df$yoe_total <- factor(df$yoe_total, yoe_levels, ordered = TRUE)
df$yoe_in_field <- factor(df$yoe_in_field, yoe_levels, ordered = TRUE)

#' Check if a given row of input has valid age & work experience
#' @param age age in a row of data
#' @param yoe_total total year of experience in a row of data
#' @param yoe_in_field year of experience in field in a row of data
#' @return A logical value indicate whether the row of data is valid
isValid <- function(age, yoe_total, yoe_in_field) {
  if (age == "under 18") {
    return(FALSE)
  }
  if (yoe_in_field > yoe_total) {
    return(FALSE)
  }
  if (age == "18-24" & yoe_total > "5-7 years") {
    return(FALSE)
  }
  if (age == "25-34" & yoe_total > "11 - 20 years") {
    return(FALSE)
  }
  if (age == "35-44" & yoe_total > "21 - 30 years") {
    return(FALSE)
  }
  if (age == "45-54" & yoe_total > "31 - 40 years") {
    return(FALSE)
  }
  return(TRUE)
}

res <- mapply(isValid, df$age, df$yoe_total, df$yoe_in_field)
print(paste('Obs before:', dim(df)[1]))

```

```
[1] "Obs before: 23374"
```

```

df <- df[res,]
print(paste('Obs after: ', dim(df)[1]))

```

```
[1] "Obs after: 23116"
```

- e. I will be using the IQR method to identify outliers. (see citation) This method might remove some realistic observations, but it will help stabilize the distribution of salary. This may help with modeling as some models are very sensitive to outliers. As shown by the histogram, the filtered salary is somewhat normally distributed, which can be a desirable property for certain types of model.

```
q <- quantile(df$salary, c(0.25, 0.75))
iqr <- IQR(df$salary)
min_val <- q[1] - 1.5 * iqr
max_val <- q[2] + 1.5 * iqr


print(paste('Obs before:', dim(df)[1]))
```

```
[1] "Obs before: 23116"
```

```
df <- subset(df, salary >= min_val & salary <= max_val)
print(paste('Obs after: ', dim(df)[1]))
```

```
[1] "Obs after: 22207"
```

```
hist(df$salary)
```



Palindromic Numbers

a. Here's the function.

```
#' Check if a number is palindromic
#' @param n Positive integer
#' @return A list with two elements:
#'         - isPalindromic: Boolean indicating if the input is palindromic
#'         - reversed: The input with its digits reversed
isPalindromic <- function(n) {
  # check invalid input
  if (!is.numeric(n) | !(all.equal(n, as.integer(n)) == TRUE) | n <= 0) {
    stop('Invalid input: positive integer only')
  }

  # convert to vector of character
  s <- strsplit(as.character(n), '')[[1]]
  # reverse the vector and convert back to number
  reverse_n <- as.numeric(paste(s[length(s):1], collapse = ''))

  return(list(isPalindromic=(all.equal(n, reverse_n) == TRUE), reversed=reverse_n))
}
```

b. Here's the function.

```
#' Get next palindromic number that's strictly greater than input
#' @param n Positive integer
#' @return Next palindromic number greater than the input
nextPalindrome <- function(n) {
  # check invalid input
  if (!is.numeric(n) | !(all.equal(n, as.integer(n)) == TRUE) | n <= 0) {
    stop('Invalid input: positive integer only')
  }

  next_n <- n + 1
  while (!isPalindromic(next_n)$isPalindromic) {
    next_n <- next_n + 1
  }

  return(next_n)
}
```

c. Here are the results.

```
# i
nextPalindrome(391)
```

```
[1] 393
```

```
# ii
nextPalindrome(9928)
```

```
[1] 9999
```

```
# iii
nextPalindrome(19272719)
```

```
[1] 19277291
```

```
# iv
nextPalindrome(109)
```

```
[1] 111
```

```
# v
nextPalindrome(2)
```

```
[1] 3
```

Citation & Link to GitHub

- [Use of ordered factor & mapply\(\) in Q2d](#)
- [Outlier detection in Q2e](#)
- [GitHub Repo of this Pset](#)